CS553 Homework #5

Sort on Single Shared Memory Node

Bhavya Chawla (bchawla@hawk.iit.edu) A2051695

Introduction

The goal of this programming assignment is to enable you to gain experience programming with external data sort and multi-threaded programming.

Table 1: Performance evaluation of Single Node TeraSort (using best # of threads for each case)

Experiment	Shared Memory (1GB)	Linux Sort (1GB)	Shared Memory (4GB)	Linux Sort (4GB)	Shared Memory (16GB)	Linux Sort (16GB)	Shared Memory (64GB)	Linux Sort (64GB)
Number of Threads	2	2	2	2	2	2	2	2
Sort Approach (e.g. in-memory / external)	In-Memory	In-Memory	External	External	External	External	External	External
Sort Algorithm (e.g. quicksort / mergesort / etc)	Mergesort	MergeSort	Mergesort	MergeSort	Mergesort	Mergesort	Mergesort	Mergesort
Data Read (GB)	1	1	8	8	32	32	128	128
Data Write (GB)	1	1	8	8	32	32	128	128
Sort Time (sec)	18.79	19.387	295.56	118	1945.78	523	6733.52	2306
Overall I/O Throughput (MB/sec)	106.43	103.16	54.13	135.59	32.89	122.37	38.01	111.01
Overall CPU Utilization (%)	37.45	27.09	43.52	31.95	48.83	37.41	50.66	42.26
Average Memory Utilization (GB)	2.9	2.1	3.4	2.56	3.9	2.99	4.0	3.3

For Linux Sort, Used **SAR** cmd for %memused - Avg memory used, and **LC_COLLATE=C** for sorting special characters as otherwise Sort cmd does not work with special characters.

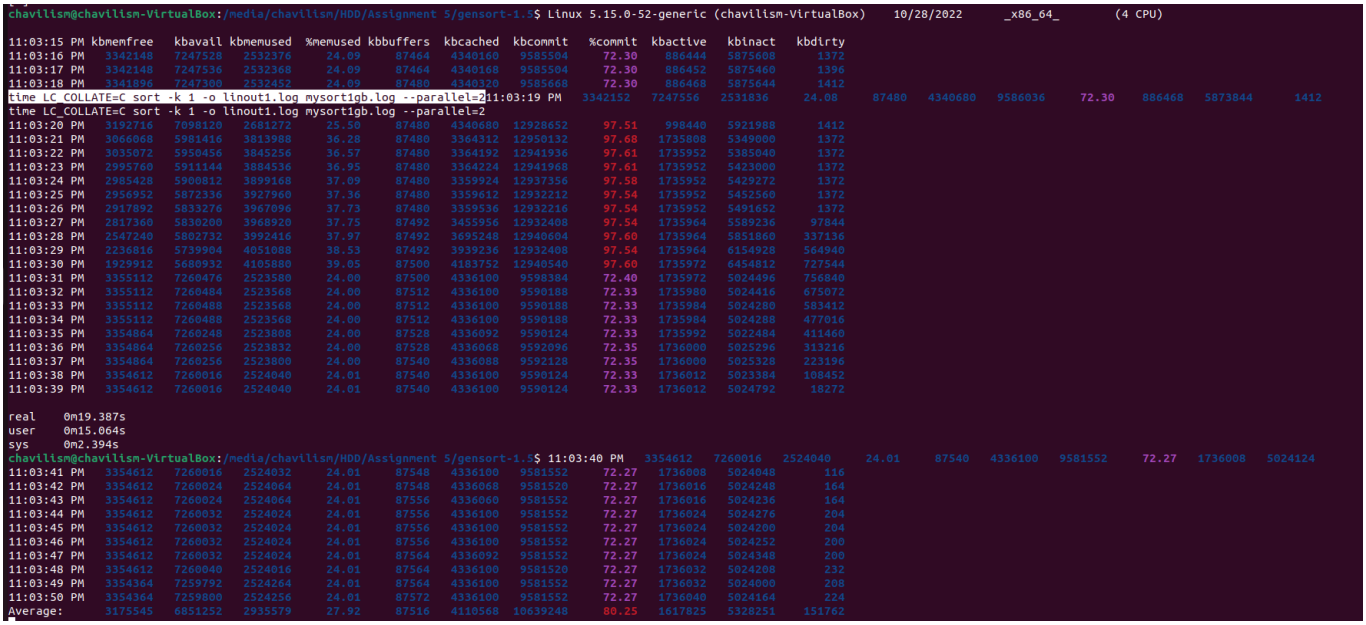

using valsort to check for order, in the output files,

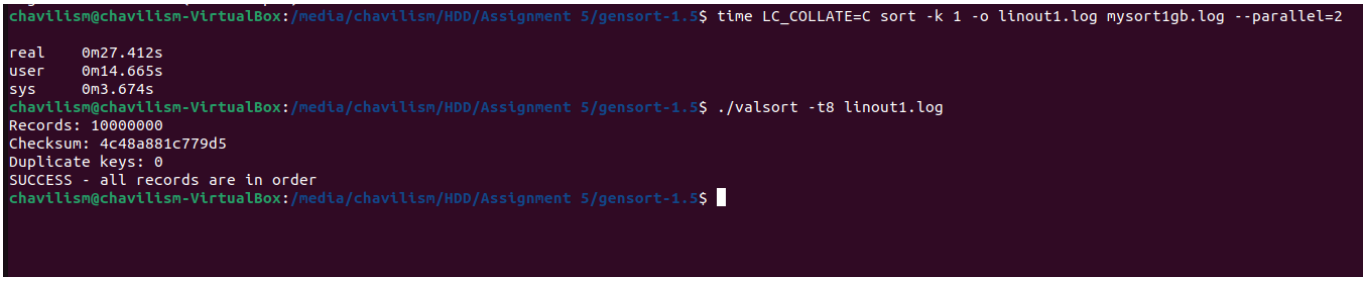

Calculations:

1. Overall I/O is calculated by (Data Read(MB) + Data Write(MB))/Sorting time(s)
2. Average Memory Utilization is calculated by considering 8GB (8000MB) memory allocated for the virtual machine and then per the %memused/CPU Utilization,
   i.e 37.45% of 8000MB
3. Data Read and Data Write is depended upon sort type, In-memory or External sort,
   for external sort per the algorithm, we are reading and writing twice the amount of data to/from storage.

   For external sort k-way Merge Sort is better than other sorting techniques as with the avg. complexity O(nlgn)

1. What you will submit

2.

The grading will be done according to the rubric below:

- Shared memory sort implementation/scripts: 50 points
- README.md: 5 points
- Performance evaluation, data, explanations, etc: 40 points
- Followed instructions on deliverables: 5 points

files:

- Makefile
- mysort.c
- hw5_report.pdf
- README.md
- mysort1GB.log
- mysort4GB.log
- mysort16GB.log
- mysort64GB.log
- linsort1GB.log
- linsort4GB.log
- linsort16GB.log
- linsort64GB.log